

COMP 5313- ARTIFICIAL INTELLIGENCE

ASSIGNMENT 2

Data Mining via Dimensionality Reduction

Methodology:


I created three .py file in order to do Data Mining via Dimensionality Reduction.

I used three methods that are PCA, t-SNE and UMAP to do dimensionality reduction.

Dataset:

Here we used MNIST Fashion dataset. Fashion-MNIST is a dataset of Zalando's article images—consisting of a training set of 60,000 examples and a test set of 10,000 examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes.

This is the dataset.



Dimensionality Reduction:

Dimensionality Reduction is a powerful technique that is widely used in data analytics, machine learning and data science to help visualize data, select good features, and to train models efficiently. We use dimensionality reduction to take higher-dimensional data and represent it in a lower dimension.

PCA:

Principal component analysis, or PCA, is a dimensionality reduction method that is often used to reduce the dimensionality of large data sets, by transforming a large set of variables into a smaller one that still contains most of the information in the large set.

First, I flattened the image into one dimensional array then normalized and reshaped it before applying PCA.

```
# Creating DataFrame for train and test data
train = pd.DataFrame(x_train_pca)
train['label'] = y_train

test = pd.DataFrame(x_test_pca)
test['label'] = y_test

# Print shapes of the datasets
print("Train DataFrame shape:", train.shape)
print("Test DataFrame shape:", test.shape)
```

Train DataFrame shape: (60000, 51)

Test DataFrame shape: (10000, 51)

This is the dataframe shape after applying PCA.

```
# Displaying the first few rows of train DataFrame
print(train.head())
```

	0	1	2	3	4	5	\
0	-123.993791	1633.074396	-1211.041191	240.793118	-3.348351	-404.340455	
1	1407.928853	-451.641336	-261.027034	366.436695	215.437558	1269.183187	
2	-725.910795	-1101.838138	106.154242	210.031701	-105.123019	-53.417242	
3	31.398664	-981.067672	202.580930	378.274376	16.283660	184.904390	
4	804.119258	-1201.168720	-744.377121	-269.630116	404.982684	-150.401060	


	6	7	8	9	...	41	\
0	-91.505515	201.375258	-32.915772	-29.809373	...	-86.447730	
1	-148.350092	-224.292458	-115.631090	-229.845293	...	81.263226	
2	-2.085852	51.304638	-91.181274	-83.071415	...	126.686722	
3	-112.847785	15.280460	-344.278935	69.950293	...	-80.660449	
4	230.429128	141.440010	14.652716	-164.896822	...	67.603761	

	42	43	44	45	46	47	\
0	71.957822	-111.994070	-58.621685	-169.638475	124.868816	-29.768906	
1	57.319452	113.628641	30.954408	95.326338	33.479146	-55.020765	
2	56.953967	-72.160244	-64.945320	16.619006	45.028633	-14.221389	
3	-110.426228	-178.700780	3.253866	7.188954	-26.138708	-58.071053	
4	30.665631	124.556751	-16.643665	-23.944379	2.526679	-5.662108	

	48	49	label
0	-83.062082	50.455423	9
1	53.708005	-57.874861	0
2	26.472528	28.104035	0
3	-6.223232	-61.465319	3
4	119.565354	-24.864397	0

[5 rows x 51 columns]

Dataframe of few rows after PCA.



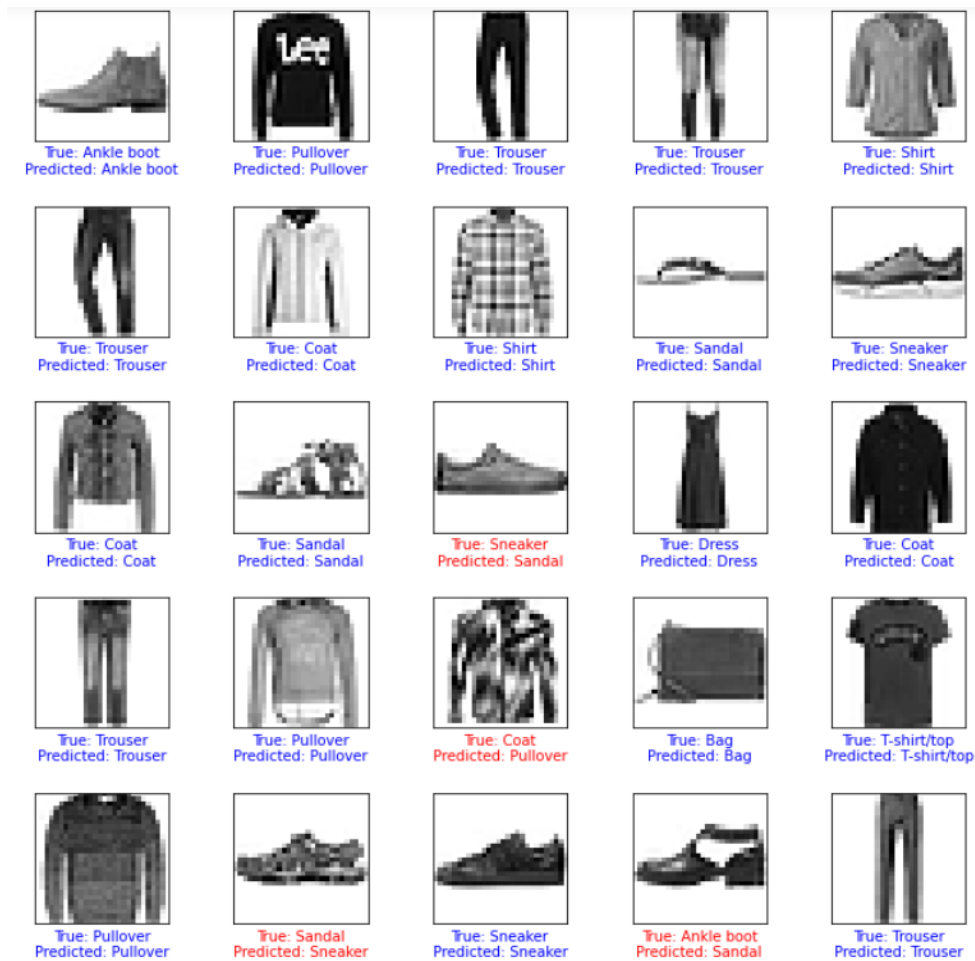
Data Visualization after PCA.

Then I applied a simple SVM classifier to classify the MNIST fashion dataset.

```
# Evaluate the performance
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.86	0.84	1000
1	0.99	0.96	0.98	1000
2	0.78	0.79	0.79	1000
3	0.87	0.89	0.88	1000
4	0.79	0.81	0.80	1000
5	0.97	0.94	0.95	1000
6	0.69	0.65	0.67	1000
7	0.92	0.94	0.93	1000
8	0.97	0.97	0.97	1000
9	0.95	0.95	0.95	1000
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000



For, this I got an accuracy of 88%.

t-SNE:

t-SNE (t-distributed Stochastic Neighbour Embedding) is an unsupervised non-linear dimensionality reduction technique for data exploration and visualizing high-dimensional data. Non-linear dimensionality reduction means that the algorithm allows us to separate data that cannot be separated by a straight line.

Similarly to PCA, I flattened the image into one dimensional array then normalized and reshaped it before applying t-SNE.

```
# Create DataFrame for train and test data
train = pd.DataFrame(x_train_tsne)
train['label'] = y_train

test = pd.DataFrame(x_test_tsne)
test['label'] = y_test

# Print shapes of the datasets
print("Train DataFrame shape:", train.shape)
print("Test DataFrame shape:", test.shape)
```

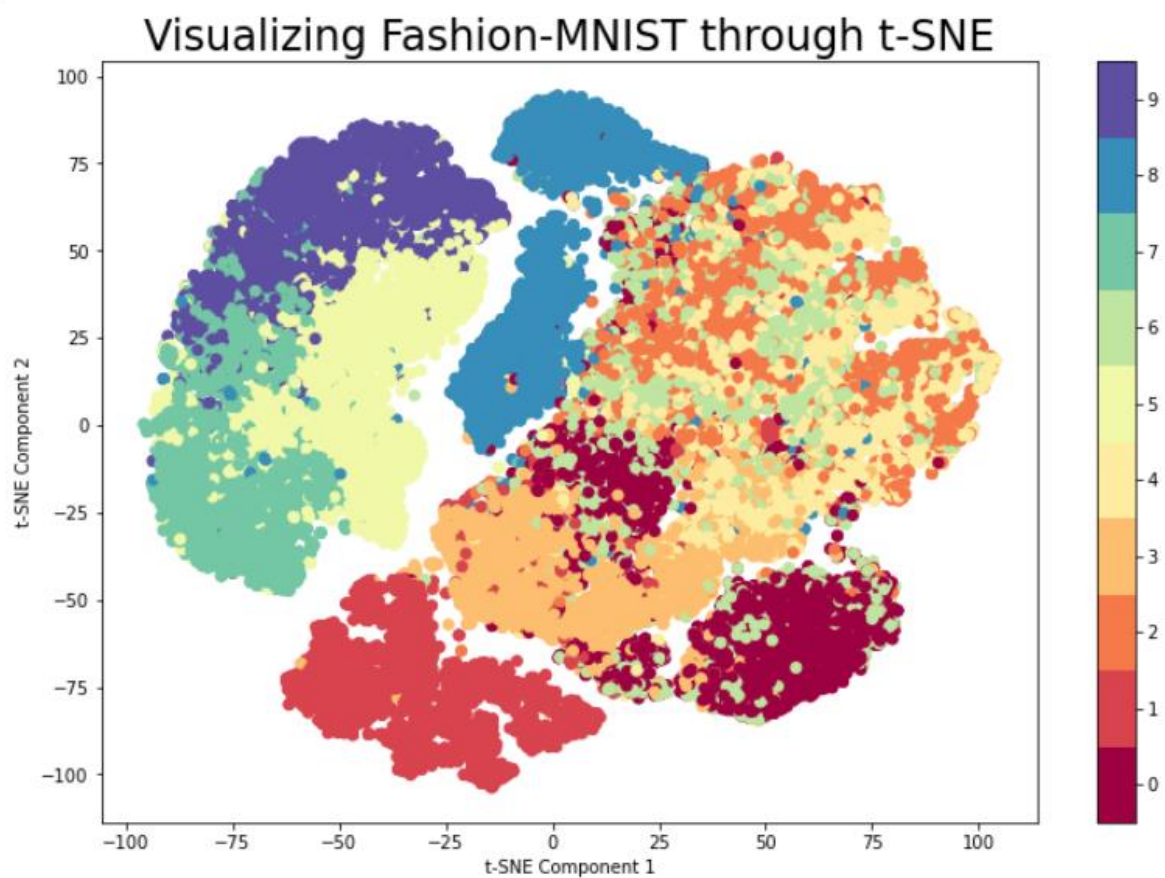
```
Train DataFrame shape: (60000, 3)
Test DataFrame shape: (10000, 3)
```

This is the dataframe shape after applying t-SNE.

```
# Display the first few rows of train DataFrame
print(train.head())
```

	0	1	label
0	-42.982185	73.936348	9
1	65.209343	-73.534035	0
2	-1.125311	-26.331484	0
3	10.212543	-31.360268	3
4	5.080834	-65.615913	0

Dataframe of few rows after t-SNE.



Data Visualization after t-SNE.

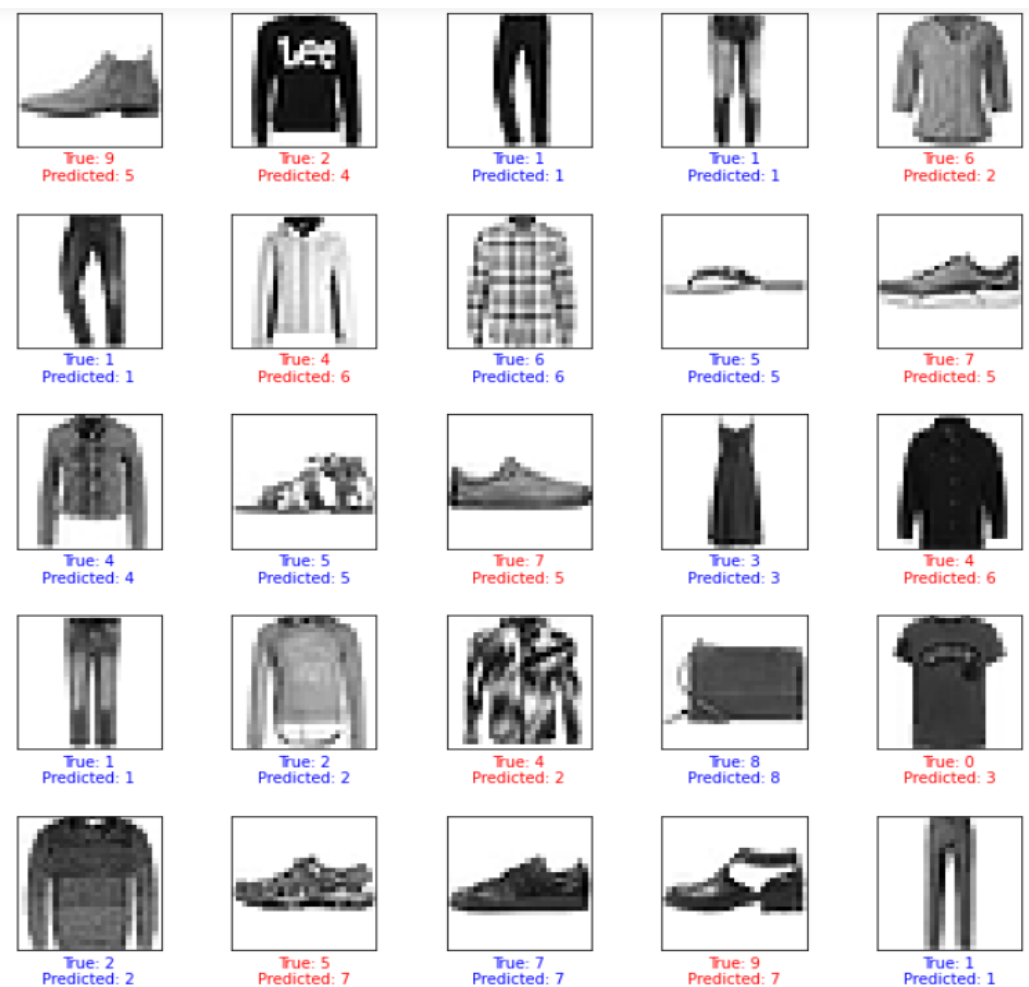
Then I applied a simple SVM classifier to classify the MNIST fashion dataset.

```
# Evaluate the performance
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0           0.50       0.49      0.49       1000
     1           1.00       0.59      0.74       1000
     2           0.41       0.44      0.42       1000
     3           0.41       0.69      0.52       1000
     4           0.48       0.44      0.46       1000
     5           0.50       0.81      0.62       1000
     6           0.33       0.27      0.29       1000
     7           0.86       0.68      0.76       1000
     8           0.80       0.84      0.82       1000
     9           0.92       0.46      0.61       1000

 accuracy              0.57      10000
 macro avg           0.62       0.57      0.57      10000
 weighted avg       0.62       0.57      0.57      10000
```



For, this I got an accuracy of 57%.

UMAP:

UMAP is a new technique by McInnes et al. Uniform Manifold Approximation and Projection (UMAP) is a dimension reduction technique that can be used for visualisation similarly to t-SNE, but also for general non-linear dimension reduction. The algorithm is founded on three assumptions about the data

1. The data is uniformly distributed on Riemannian manifold;
2. The Riemannian metric is locally constant
3. The manifold is locally connected.

Here also, I flattened the image into one dimensional array then normalized and reshaped it before applying UMAP.

```
# Create DataFrame for train and test data
train = pd.DataFrame(x_train_umap)
train['label'] = y_train

test = pd.DataFrame(x_test_umap)
test['label'] = y_test

# Print shapes of the datasets
print("Train DataFrame shape:", train.shape)
print("Test DataFrame shape:", test.shape)
```

```
Train DataFrame shape: (60000, 51)
Test DataFrame shape: (10000, 51)
```

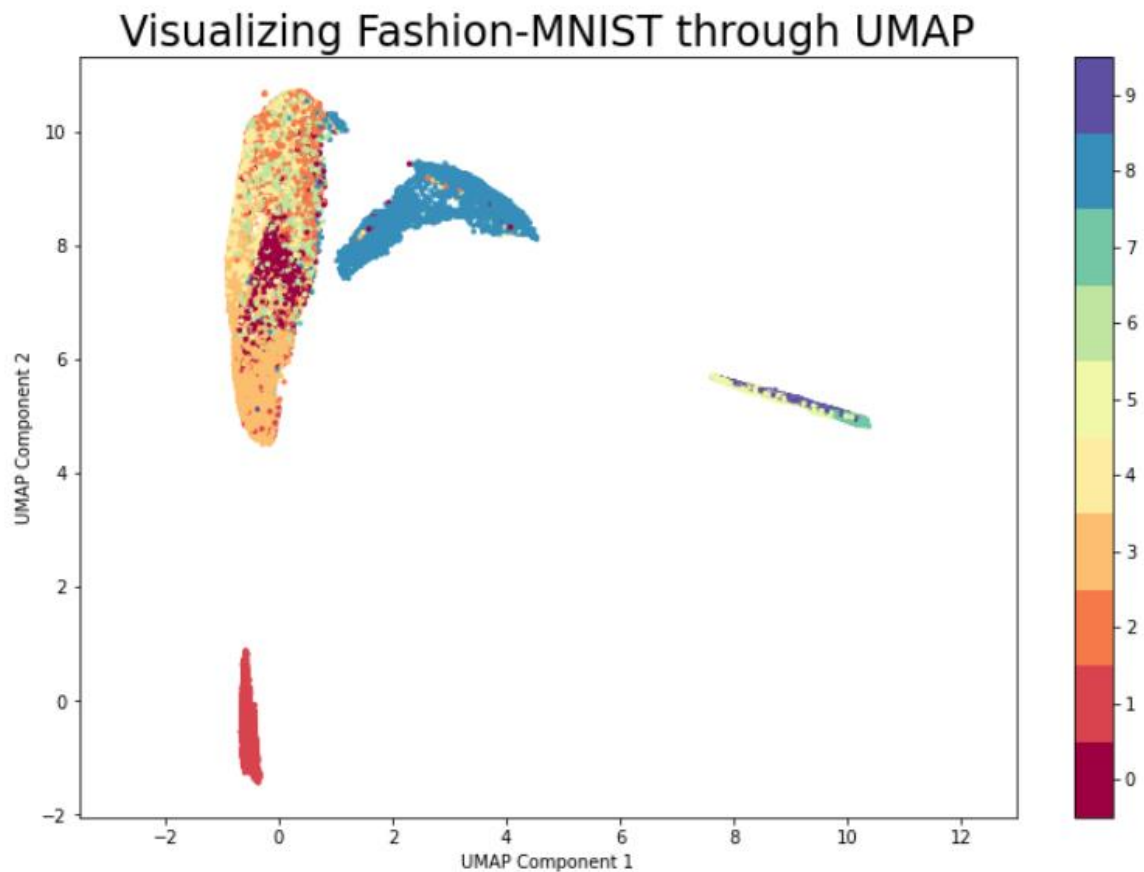
This is the dataframe shape after applying UMAP.

```
# Display the first few rows of train DataFrame
print(train.head())
```

	0	1	2	3	4	5	6	\
0	0.930203	4.629578	7.567747	4.337176	3.303705	4.401646	4.592235	
1	10.143060	1.798088	4.262380	1.502978	6.882712	5.672943	3.115500	
2	10.004306	3.413370	4.453558	3.040205	5.444238	3.964400	5.339095	
3	10.217864	3.198436	4.393587	2.458318	5.913263	4.187828	5.002822	
4	10.717884	4.024543	4.377642	1.729188	6.349250	3.768091	5.717562	
	7	8	9	...	41	42	43	44 \
0	4.509768	4.627343	3.981642	...	3.928945	4.979936	3.985417	6.529401
1	5.206645	3.885250	4.095042	...	3.880199	4.964187	3.980197	6.494970
2	5.331816	3.906265	4.716318	...	3.907916	4.943336	3.979457	6.492531
3	5.237074	4.058956	4.491087	...	3.911539	4.941634	3.982293	6.477735
4	5.038370	4.590066	4.229511	...	3.911373	4.919974	3.985278	6.456738
	45	46	47	48	49	label		
0	6.121536	3.751657	5.648276	6.366957	4.141622	9		
1	6.115186	3.736057	5.620748	6.301548	4.133403	0		
2	6.148813	3.759392	5.651603	6.372872	4.163128	0		
3	6.150653	3.753586	5.656752	6.365251	4.161204	3		
4	6.158533	3.755945	5.658025	6.366473	4.165531	0		

```
[5 rows x 51 columns]
```

Dataframe of few rows after UMAP.



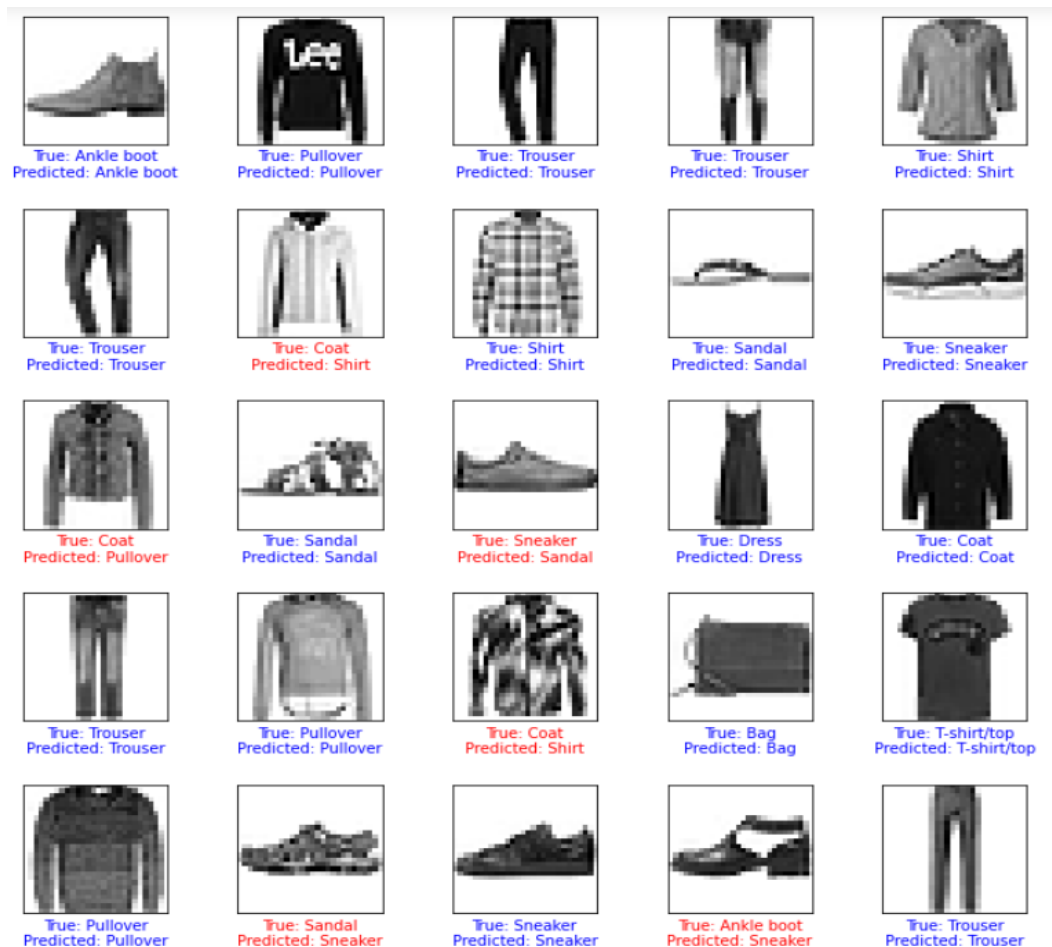
Data Visualization after UMAP.

Then I applied a simple SVM classifier to classify the MNIST fashion dataset.

```
# Evaluate the performance
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.67	0.87	0.76	1000
1	0.99	0.94	0.97	1000
2	0.56	0.62	0.59	1000
3	0.82	0.83	0.82	1000
4	0.55	0.59	0.57	1000
5	0.93	0.82	0.87	1000
6	0.49	0.32	0.39	1000
7	0.85	0.88	0.86	1000
8	0.96	0.89	0.93	1000
9	0.86	0.94	0.90	1000
accuracy			0.77	10000
macro avg	0.77	0.77	0.77	10000
weighted avg	0.77	0.77	0.77	10000



For, this I got an accuracy of 77%.

Conclusion:

As you can see above, after using all the dimensionality reduction techniques. PCA got the higher accuracy of 88% while UMAP got 77% and t-SNE got 57% for this simple SVM model. Even though, PCA and t-SNE are widely used techniques their performance suffers with large datasets and using it correctly can be challenging. While UMAP offers a number of advantages over PCA and t-SNE.

References:

- [1] <https://builtin.com/data-science/step-step-explanation-principal-component-analysis>
- [2] <https://www.kaggle.com/code/parulpandey/part1-visualizing-kannada-mnist-with-pca?scriptVersionId=29322090>
- [3] <https://www.datacamp.com/tutorial/introduction-t-sne>
- [4] <https://www.kaggle.com/code/parulpandey/visualizing-kannada-mnist-with-t-sne>
- [5] <https://umap-learn.readthedocs.io/en/latest/>
- [6] https://umap-learn.readthedocs.io/en/latest/auto_examples/plot_mnist_example.html
- [7] <https://pair-code.github.io/understanding-umap/>